

# Backfitting

DJM

11 March 2017

## Backfitting

```
backfitlm <- function(y, X, max.iter=100, small.enough=1e-4, track=FALSE){  
  # This function performs linear regression by "backfitting"  
  # Inputs: y - the response  
  #         X - the design matrix  
  #         max.iter - the maximum number of loops through the predictors (default to 100)  
  #         small.enough - if the mse changes by less than this, terminate (default to 1e-6)  
  X = as.matrix(X)  
  p = ncol(X) # how many covariates?  
  n = nrow(X) # how many observations  
  betas = double(p) # create a vector for our estimated coefficients (all zeros now)  
  preds = matrix(0, n, p) # a matrix to hold the partial predictions of each covariate  
  pres = matrix(y, n, p) # partial residuals begin as y (since we are predicting with 0)  
  iter = 0 # initialize our iteration  
  mse = mean(y^2) # initialize the MSE to SSTo  
  conv = FALSE # initialize the convergence check to FALSE  
  while(!conv && (iter < max.iter)){ # enter the loop, check conditions (note && not &)  
    iter = iter + 1 # update the iteration count  
    for(j in 1:p){ # loop over all predictors  
      pres[,j] = y - rowSums(preds[, -j]) # partial residuals (ignoring current predictor)  
      ## same as X[, -j] %*% betas[-j], same as apply(preds[, -j], 1, sum)  
      mod = lm(pres[,j] ~ X[,j]-1) # regress current predictor on partial residuals, no intercept, if g  
      ## mod = npreg(pres[,j]~X[,j], tol=1e-4, ftol=1e-4) # if we want a gam instead  
      betas[j] = coefficients(mod) # get out the single coefficient, if gam, remove this line  
      preds[,j] = fitted(mod) # update the predictions from this column  
    }  
    msenew = sqrt(mean((y - rowSums(preds)))^2) # get the updated MSE after a pass  
    conv = (abs(mse-msenew)<small.enough) # check how different our MSE was from previous  
    mse = msenew # save the new MSE  
    if(track) cat(iter/max.iter, " mse = ", mse, "\n")  
  }  
  return(list(bhat=betas, pres=pres, preds = preds, mse=mse)) # return our coefficients  
}
```

- Generate a design matrix  $X$  with 100 observations and  $p = 10$  covariates and a response variable  $y$  using a linear model. Test the (now complete) `backfitlm` on this data and compare the results to `lm`. Should there be an intercept in either version?

```
set.seed(03-09-2017)  
n = 100  
p = 10  
X = matrix(runif(n*p,-1,1), n)  
b = 10:1 # true betas  
y = X %*% b + rnorm(n, sd=.5)  
bhat.lm = coef(lm(y~X-1)) # no intercept  
bhat.bf = backfitlm(y,X)$bhat # also no intercept
```

```
round(rbind(bhat.lm,bhat.bf),3)
```

```
##           X1      X2      X3      X4      X5      X6      X7      X8      X9     X10
## bhat.lm 10.096  8.974  7.914  6.927  5.952  4.872  4.014  2.984  1.765  0.995
## bhat.bf 10.096  8.974  7.914  6.927  5.952  4.872  4.014  2.984  1.765  0.995
```

Notice that the estimated coefficients are exactly the same. I didn't generate data with an intercept, so I didn't let `lm` estimate one. You could have though. You just need to include a column of ones in the `X` matrix.